

Text analysis

You can download this .qmd file from [here](#). Just hit the Download Raw File button.

We will build on techniques you learned in SDS 164 using parts of [Text Mining with R](#) by Silge and Robinson.

Text analysis of books from Project Gutenberg

We will use the `gutenbergr` package to obtain several works from [Project Gutenberg](#) to examine using text analysis tools.

```
# How I obtained the three works from Project Gutenberg

# Notes:
# - might have to find mirror at https://www.gutenberg.org/MIRRORS.ALL
# - 84 = Frankenstein; 345 = Dracula; 43 = Jekyll and Hyde

# three_works <- gutenberg_download(
#   c(84, 345, 43),
#   meta_fields = "title",
#   mirror = "http://mirror.csclub.uwaterloo.ca/gutenberg/")

# write_csv(three_works, "~/264_fall_2024/Data/three_works.csv")

# three_works <- read_csv("https://proback.github.io/264_fall_2024/Data/three_works.csv")
# three_works2 <- read_csv("Data/three_works.csv")

library(RCurl)
```

Attaching package: 'RCurl'

The following object is masked from 'package:tidyr':

complete

```
three_works <- read_csv(  
  file = getURL("https://raw.githubusercontent.com/proback/264_fall_2024/refs/heads/main/D
```

Rows: 25399 Columns: 3

-- Column specification -----

Delimiter: ","

chr (2): text, title

dbl (1): gutenber_id

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
three_works |> count(title)
```

A tibble: 3 x 2

	title	n
	<chr>	<int>
1	Dracula	15491
2	Frankenstein; Or, The Modern Prometheus	7357
3	The Strange Case of Dr. Jekyll and Mr. Hyde	2551

```
three_works
```

A tibble: 25,399 x 3

	gutenber_id	text	title
	<dbl>	<chr>	<chr>
1	43	The Strange Case Of Dr. Jekyll And Mr. Hyde	The Strange Case of~
2	43	<NA>	The Strange Case of~
3	43	by Robert Louis Stevenson	The Strange Case of~
4	43	<NA>	The Strange Case of~
5	43	<NA>	The Strange Case of~
6	43	Contents	The Strange Case of~
7	43	<NA>	The Strange Case of~

```

8           43 <NA>           The Strange Case of~
9           43 STORY OF THE DOOR   The Strange Case of~
10          43 <NA>           The Strange Case of~
# i 25,389 more rows

```

```

frankenstein <- three_works |>
  filter(str_detect(title, "Frankenstein"))

```

We will begin by looking at a single book (Frankenstein) and then we'll compare and contrast 3 books (Frankenstein, Dracula, and Jekyll and Hyde).

1. Tidy Text Mining!

Now it's time to tokenize and tidy this text data.

```

tidy_book <- frankenstein |>
  mutate(line = row_number()) |>
  unnest_tokens(word, text, token = "words") # (new name, input)
# default for unnest_tokens is token = "words", but can also use
# "characters", "ngrams" with say n=2, "sentences", "regex" with
# pattern, "tweets" with strip_url, etc.

```

```

tidy_book # one row per word, instead of one per line

```

```

# A tibble: 76,204 x 4

```

	gutenberg_id	title	line	word
	<dbl>	<chr>	<int>	<chr>
1	84	Frankenstein; Or, The Modern Prometheus	1	frankenstein
2	84	Frankenstein; Or, The Modern Prometheus	2	<NA>
3	84	Frankenstein; Or, The Modern Prometheus	3	or
4	84	Frankenstein; Or, The Modern Prometheus	3	the
5	84	Frankenstein; Or, The Modern Prometheus	3	modern
6	84	Frankenstein; Or, The Modern Prometheus	3	prometheus
7	84	Frankenstein; Or, The Modern Prometheus	4	<NA>
8	84	Frankenstein; Or, The Modern Prometheus	5	by
9	84	Frankenstein; Or, The Modern Prometheus	5	mary
10	84	Frankenstein; Or, The Modern Prometheus	5	wollstonecraft

```

# i 76,194 more rows

```

```
frankenstein |> slice_tail(n = 10)
```

```
# A tibble: 10 x 3
```

	gutenberg_id	text	title
	<dbl>	<chr>	<chr>
1	84	shall die, and what I now feel be no longer felt. Soon th~	Fran~
2	84	miseries will be extinct. I shall ascend my funeral pile ~	Fran~
3	84	exult in the agony of the torturing flames. The light of ~	Fran~
4	84	will fade away; my ashes will be swept into the sea by th~	Fran~
5	84	will sleep in peace, or if it thinks, it will not surely ~	Fran~
6	84	Farewell."	Fran~
7	84	<NA>	Fran~
8	84	He sprang from the cabin-window as he said this, upon the~	Fran~
9	84	which lay close to the vessel. He was soon borne away by ~	Fran~
10	84	lost in darkness and distance.	Fran~

```
tidy_book |> slice_tail(n = 20)
```

```
# A tibble: 20 x 4
```

	gutenberg_id	title	line	word
	<dbl>	<chr>	<int>	<chr>
1	84	Frankenstein; Or, The Modern Prometheus	7356	which
2	84	Frankenstein; Or, The Modern Prometheus	7356	lay
3	84	Frankenstein; Or, The Modern Prometheus	7356	close
4	84	Frankenstein; Or, The Modern Prometheus	7356	to
5	84	Frankenstein; Or, The Modern Prometheus	7356	the
6	84	Frankenstein; Or, The Modern Prometheus	7356	vessel
7	84	Frankenstein; Or, The Modern Prometheus	7356	he
8	84	Frankenstein; Or, The Modern Prometheus	7356	was
9	84	Frankenstein; Or, The Modern Prometheus	7356	soon
10	84	Frankenstein; Or, The Modern Prometheus	7356	borne
11	84	Frankenstein; Or, The Modern Prometheus	7356	away
12	84	Frankenstein; Or, The Modern Prometheus	7356	by
13	84	Frankenstein; Or, The Modern Prometheus	7356	the
14	84	Frankenstein; Or, The Modern Prometheus	7356	waves
15	84	Frankenstein; Or, The Modern Prometheus	7356	and
16	84	Frankenstein; Or, The Modern Prometheus	7357	lost
17	84	Frankenstein; Or, The Modern Prometheus	7357	in
18	84	Frankenstein; Or, The Modern Prometheus	7357	darkness
19	84	Frankenstein; Or, The Modern Prometheus	7357	and
20	84	Frankenstein; Or, The Modern Prometheus	7357	distance

What are the most common words?

```
tidy_book |>
  count(word, sort = TRUE)
```

```
# A tibble: 7,077 x 2
  word      n
  <chr> <int>
1 the    4195
2 and    2976
3 i      2846
4 of     2642
5 to     2089
6 my     1776
7 a      1391
8 in     1128
9 was    1021
10 that  1017
# i 7,067 more rows
```

Stop words (get rid of common but not useful words)

Note: If you get “Error in loadNamespace(name) : there is no package called ‘stopwords’”, first install package `stopwords`.

```
get_stopwords() |> print(n = 50) # snowball is default - somewhat smaller
```

```
# A tibble: 175 x 2
  word      lexicon
  <chr>    <chr>
1 i       snowball
2 me      snowball
3 my      snowball
4 myself  snowball
5 we      snowball
6 our     snowball
7 ours    snowball
8 ourselves snowball
9 you     snowball
10 your   snowball
```

11	yours	snowball
12	yourself	snowball
13	yourselves	snowball
14	he	snowball
15	him	snowball
16	his	snowball
17	himself	snowball
18	she	snowball
19	her	snowball
20	hers	snowball
21	herself	snowball
22	it	snowball
23	its	snowball
24	itself	snowball
25	they	snowball
26	them	snowball
27	their	snowball
28	theirs	snowball
29	themselves	snowball
30	what	snowball
31	which	snowball
32	who	snowball
33	whom	snowball
34	this	snowball
35	that	snowball
36	these	snowball
37	those	snowball
38	am	snowball
39	is	snowball
40	are	snowball
41	was	snowball
42	were	snowball
43	be	snowball
44	been	snowball
45	being	snowball
46	have	snowball
47	has	snowball
48	had	snowball
49	having	snowball
50	do	snowball
# i 125 more rows		

```
get_stopwords(source = "smart") |> print(n = 50)
```

```
# A tibble: 571 x 2
  word      lexicon
  <chr>    <chr>
1 a       smart
2 a's     smart
3 able    smart
4 about   smart
5 above   smart
6 according smart
7 accordingly smart
8 across  smart
9 actually smart
10 after   smart
11 afterwards smart
12 again   smart
13 against smart
14 ain't   smart
15 all     smart
16 allow   smart
17 allows  smart
18 almost  smart
19 alone   smart
20 along   smart
21 already smart
22 also    smart
23 although smart
24 always  smart
25 am      smart
26 among   smart
27 amongst smart
28 an      smart
29 and     smart
30 another smart
31 any     smart
32 anybody smart
33 anyhow  smart
34 anyone  smart
35 anything smart
36 anyway  smart
37 anyways smart
```

```
38 anywhere    smart
39 apart       smart
40 appear      smart
41 appreciate  smart
42 appropriate smart
43 are         smart
44 aren't      smart
45 around      smart
46 as          smart
47 aside       smart
48 ask         smart
49 asking      smart
50 associated  smart
# i 521 more rows
```


```
# will sometimes want to store if using over and over
# - later with shiny apps will have to store and write as data file
smart_stopwords <- get_stopwords(source = "smart")
```

Try out using different languages (`language`) and different lexicons (`source`).

Another try at most common words

```
tidy_book |>
  anti_join(smart_stopwords) |>
  count(word, sort = TRUE) |>
  filter(word != "NA") |>
  slice_max(n, n = 20) |>
  ggplot(aes(fct_reorder(word, n), n)) +
  geom_col() +
  coord_flip()
```

Joining with ``by = join_by(word)``



Sentiment analysis

Explore some sentiment lexicons. You'll want to match your choice of sentiment lexicon to your purpose:

- `afinn`: scored from -5 (very negative) to +5 (very positive)
- `nrc`: words are labeled with emotions like anger, fear, sadness, etc. There can be more than one row per word.
- `bing`: binary - listed words are either negative or positive

```
get_sentiments(lexicon = "afinn")
```

A tibble: 2,477 x 2

	word	value
	<chr>	<dbl>
1	abandon	-2
2	abandoned	-2
3	abandons	-2
4	abducted	-2
5	abduction	-2
6	abductions	-2
7	abhor	-3

```

8 abhorred      -3
9 abhorrent     -3
10 abhors       -3
# i 2,467 more rows

```

```
get_sentiments(lexicon = "nrc")
```

```

# A tibble: 13,872 x 2
  word      sentiment
  <chr>     <chr>
1 abacus    trust
2 abandon   fear
3 abandon   negative
4 abandon   sadness
5 abandoned anger
6 abandoned fear
7 abandoned negative
8 abandoned sadness
9 abandonment anger
10 abandonment fear
# i 13,862 more rows

```

```
get_sentiments(lexicon = "bing")
```

```

# A tibble: 6,786 x 2
  word      sentiment
  <chr>     <chr>
1 2-faces   negative
2 abnormal  negative
3 abolish   negative
4 abominable negative
5 abominably negative
6 abominate negative
7 abomination negative
8 abort     negative
9 aborted   negative
10 aborts    negative
# i 6,776 more rows

```

```

afinn_sentiments <- get_sentiments(lexicon = "afinn")
nrc_sentiments <- get_sentiments(lexicon = "nrc")
bing_sentiments <- get_sentiments(lexicon = "bing")

```

Implement sentiment analysis using an `inner_join()`, so you only consider words both in your text and in the lexicon.

```

tidy_book |>
  inner_join(bing_sentiments) |>
  count(sentiment)

```

Joining with ``by = join_by(word)``

```

# A tibble: 2 x 2
  sentiment      n
  <chr>      <int>
1 negative   3742
2 positive   2983

```


What words contribute the most to sentiment scores for Frankenstein? Let's walk through this pipe step-by-step.

```

tidy_book |>
  inner_join(bing_sentiments) |>
  count(sentiment, word, sort = TRUE) |>
  group_by(sentiment) |>
  slice_max(n, n = 10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(word, n), y = n, fill = sentiment)) +
    geom_col() +
    coord_flip() +
    facet_wrap(~ sentiment, scales = "free")

```

Joining with ``by = join_by(word)``



```
# Check out which words are associated with specific nrc emotions
get_sentiments("nrc") |>
  count(sentiment)
```

```
# A tibble: 10 x 2
  sentiment      n
  <chr>      <int>
1 anger      1245
2 anticipation 837
3 disgust    1056
4 fear       1474
5 joy         687
6 negative    3316
7 positive    2308
8 sadness     1187
9 surprise     532
10 trust      1230
```

```
get_sentiments("nrc") |>
  filter(sentiment == "joy") |>
  inner_join(tidy_book) |>
```

```
count(word, sort = TRUE)
```

Joining with `by = join_by(word)`

```
# A tibble: 308 x 2
```

	word	n
	<chr>	<int>
1	found	87
2	friend	71
3	love	59
4	hope	50
5	happiness	49
6	happy	46
7	sun	45
8	joy	42
9	affection	40
10	journey	36

```
# i 298 more rows
```

```
get_sentiments("nrc") |>  
  filter(sentiment == "anger") |>  
  inner_join(tidy_book) |>  
  count(word, sort = TRUE)
```

Joining with `by = join_by(word)`

```
# A tibble: 370 x 2
```

	word	n
	<chr>	<int>
1	death	79
2	miserable	65
3	misery	54
4	words	54
5	despair	49
6	horror	45
7	fear	40
8	possessed	36
9	fiend	33
10	feeling	27

```
# i 360 more rows
```

Make a wordcloud for Frankenstein.

```
# wordcloud wants a column with words and another column with counts
words <- tidy_book |>
  anti_join(stop_words) |>
  count(word) |>
  filter(word != "NA") |>
  arrange(desc(n))

# Note: this will look better in html than in the Plots window in RStudio
wordcloud(
  words = words$word,
  freq = words$n,
  max.words = 100,
  random.order = FALSE
)
```



```
# See Z's R Tip of the Day for suggestions on options
wordcloud(
  words = words$word,
  freq = words$n,
```

```

max.words = 200,
random.order = FALSE,
rot.per = 0.35,
scale = c(3.5, 0.25),
colors = brewer.pal(6, "Dark2"))

```



```

# Or for even cooler looks, use wordcloud2 (for html documents)
words_df <- words |>
  slice_head(n = 80) |>
  data.frame()

wordcloud2(
  words_df,
  size = .25,
  shape = 'circle',
  minSize = 10
)

```




```
# A couple of helpful links for customizing wordclouds:
#   https://www.youtube.com/watch?v=0cToDzeDLRI
#   https://towardsdatascience.com/create-a-word-cloud-with-r-bde3e7422e8a
```

You could do cool stuff here, like color the words by sentiment!

2. What is a document about?

Above, we read in a corpus called `three_works`. We'll use that here!

Count the word frequencies by title in this collection.

```
book_words <- three_works |>
  group_by(title) |>
  mutate(linenumber = row_number()) |>
  ungroup() |>
  unnest_tokens(word, text)

book_word_count <- book_words |>
  count(word, title, sort = TRUE)

book_word_count
```

```
# A tibble: 20,714 x 3
  word title n
  <chr> <chr> <int>
1 the Dracula 7915
2 and Dracula 5907
3 i Dracula 4801
4 to Dracula 4666
5 the Frankenstein; Or, The Modern Prometheus 4195
6 of Dracula 3634
7 and Frankenstein; Or, The Modern Prometheus 2976
8 a Dracula 2954
9 i Frankenstein; Or, The Modern Prometheus 2846
10 of Frankenstein; Or, The Modern Prometheus 2642
# i 20,704 more rows
```

Look at positive/negative sentiment trajectory over the novels

```

book_words |>
  inner_join(bing_sentiments) |>
  count(title, index = linenumber %% 80, sentiment) |>
  # index approximates a chapter (every 80 lines)
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0) |>
  mutate(sentiment = positive - negative) |>
  ggplot(aes(x = index, y = sentiment, fill = title)) +
    geom_col(show.legend = FALSE) +
    facet_wrap(~title, ncol = 2, scales = "free_x")

```

Joining with `by = join_by(word)`



Calculate tf-idf.

The tf-idf statistic is term frequency times inverse document frequency, a quantity used for identifying terms that are especially important to a particular document. The idea of tf-idf is to find the important words for the content of each document by decreasing the weight for commonly used words and increasing the weight for words that are not used very much in a collection or corpus of documents. We want to find words that define one document as opposed to others.

- tf = term frequency = proportion of times a term appears in a document.
- idf = inverse document frequency = $\log(\text{number of documents} / \text{number of documents with the term})$, so that terms that appear in fewer documents are weighted higher, since those rarer words provide more information.

There's really no theory behind multiplying the two together - it just tends to work in practice. See [this wikipedia entry](#) for more details. (See also [this site](#) for a nice description of weaknesses of tf - idf .)

```
book_tfidf <- book_word_count |>
  bind_tf_idf(word, title, n)

book_tfidf  # note idf = 0 when it appears in every document

# A tibble: 20,714 x 6
  word title n tf idf tf_idf
  <chr> <chr> <int> <dbl> <dbl> <dbl>
1 the Dracula 7915 0.0480 0 0
2 and Dracula 5907 0.0358 0 0
3 i Dracula 4801 0.0291 0 0
4 to Dracula 4666 0.0283 0 0
5 the Frankenstein; Or, The Modern Prometheus 4195 0.0550 0 0
6 of Dracula 3634 0.0220 0 0
7 and Frankenstein; Or, The Modern Prometheus 2976 0.0391 0 0
8 a Dracula 2954 0.0179 0 0
9 i Frankenstein; Or, The Modern Prometheus 2846 0.0373 0 0
10 of Frankenstein; Or, The Modern Prometheus 2642 0.0347 0 0
# i 20,704 more rows
```

Find *high* tf - idf words. The highest words will appear relatively often in one document, but not at all in others.

```
book_tfidf |>
  arrange(-tf_idf)


# A tibble: 20,714 x 6
  word title n tf idf tf_idf
  <chr> <chr> <int> <dbl> <dbl> <dbl>
1 utterson The Strange Case of Dr. Jekyll and Mr.~ 128 0.00489 1.10 0.00537
2 jekyll The Strange Case of Dr. Jekyll and Mr.~ 84 0.00321 1.10 0.00353
```

3	poole	The Strange Case of Dr. Jekyll and Mr.~	61	0.00233	1.10	0.00256
4	van	Dracula	323	0.00196	1.10	0.00215
5	helsing	Dracula	301	0.00182	1.10	0.00200
6	hyde	The Strange Case of Dr. Jekyll and Mr.~	98	0.00375	0.405	0.00152
7	lucy	Dracula	223	0.00135	1.10	0.00148
8	mina	Dracula	210	0.00127	1.10	0.00140
9	elizabeth	Frankenstein; Or, The Modern Prometheus	88	0.00115	1.10	0.00127
10	jonathan	Dracula	181	0.00110	1.10	0.00120

i 20,704 more rows

How can we visualize this? Let's go step-by-step.

```
book_tfidf |>
  group_by(title) |>
  arrange(desc(tf_idf)) |>
  slice_max(tf_idf, n = 10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(word, tf_idf), y = tf_idf, fill = title)) +
    geom_col(show.legend = FALSE) +
    coord_flip() +
    facet_wrap(~title, scales = "free")
```



```
# kind of boring - mostly proper nouns
```

N-grams... and beyond!

Let's return to Frankenstein and look at 2-word combinations:

```
tidy_ngram <- frankenstein |>
  unnest_tokens(bigram, text, token = "ngrams", n = 2) |>
  filter(bigram != "NA")
```

```
tidy_ngram
```

```
# A tibble: 68,847 x 3
```

	gutenberg_id	title	bigram
	<dbl>	<chr>	<chr>
1	84	Frankenstein; Or, The Modern Prometheus	or the
2	84	Frankenstein; Or, The Modern Prometheus	the modern
3	84	Frankenstein; Or, The Modern Prometheus	modern prometheus
4	84	Frankenstein; Or, The Modern Prometheus	by mary
5	84	Frankenstein; Or, The Modern Prometheus	mary wollstonecraft
6	84	Frankenstein; Or, The Modern Prometheus	wollstonecraft godwin
7	84	Frankenstein; Or, The Modern Prometheus	godwin shelley
8	84	Frankenstein; Or, The Modern Prometheus	letter 1
9	84	Frankenstein; Or, The Modern Prometheus	letter 2
10	84	Frankenstein; Or, The Modern Prometheus	letter 3

```
# i 68,837 more rows
```

What are the most common bigrams?

```
tidy_ngram |>
  count(bigram, sort = TRUE)
```

```
# A tibble: 38,574 x 2
```

	bigram	n
	<chr>	<int>
1	of the	501
2	of my	264
3	in the	246
4	i was	213

```

5 i had      207
6 that i     198
7 and i      192
8 and the    182
9 to the     181
10 which i    145
# i 38,564 more rows

```

Let's use `separate()` from `tidyr` to remove stop words.

```

# stop_words contains 1149 words from 3 lexicons
bigrams_filtered <- tidy_ngram |>
  separate(bigram, c("word1", "word2"), sep = " ") |>
  filter(!word1 %in% stop_words$word,
         !word2 %in% stop_words$word) |>
  count(word1, word2, sort = TRUE)
bigrams_filtered

```

```

# A tibble: 4,677 x 3
  word1      word2      n
  <chr>      <chr>    <int>
1 natural    philosophy    11
2 dear       victor       10
3 native     country       10
4 de         lacey         9
5 fellow     creatures      8
6 poor       girl          8
7 mont       blanc         7
8 native     town          6
9 cornelius  agrippa       5
10 countenance expressed  5
# i 4,667 more rows

```

Now extend from a single document to our collection of documents. See which two-word combinations best identify books in the collection.

```

book_twowords <- three_works |>
  group_by(title) |>
  mutate(linenumber = row_number()) |>
  ungroup() |>
  unnest_tokens(bigram, text, token = "ngrams", n = 2) |>

```

```

    filter(bigram != "NA")

book_twowords |>
  count(bigram, sort = TRUE)

# A tibble: 102,837 x 2
  bigram      n
  <chr>    <int>
1 of the   1494
2 in the    952
3 to the    596
4 and the   579
5 and i     554
6 it was    526
7 that i    526
8 on the    507
9 i was     484
10 i had    461
# i 102,827 more rows

bigrams_filtered <- book_twowords |>
  separate(bigram, c("word1", "word2"), sep = " ") |>
  filter(!word1 %in% stop_words$word,
         !word2 %in% stop_words$word) |>
  count(word1, word2, sort = TRUE) |>
  filter(!is.na(word1) & !is.na(word2))

bigrams_filtered

# A tibble: 13,951 x 3
  word1    word2      n
  <chr>    <chr>    <int>
1 van      helsing   282
2 madam     mina     82
3 lord      godalming  63
4 dr        van      60
5 dr        seward    55
6 friend    john     54
7 seward's  diary    39
8 poor      dear     34

```

```

9 harker's journal      31
10 _dr      seward's    26
# i 13,941 more rows

```

```

bigrams_united <- bigrams_filtered |>
  unite(bigram, word1, word2, sep = " ")

```

```
bigrams_united
```

```
# A tibble: 13,951 x 2
```

```

  bigram      n
  <chr>    <int>
1 van helsing      282
2 madam mina       82
3 lord godalming   63
4 dr van           60
5 dr seward        55
6 friend john      54
7 seward's diary   39
8 poor dear        34
9 harker's journal  31
10 _dr seward's    26
# i 13,941 more rows

```

```

bigram_tf_idf <- book_twowords |>
  count(title, bigram) |>
  bind_tf_idf(bigram, title, n) |>
  arrange(desc(tf_idf))

```

```
bigram_tf_idf |> arrange(desc(tf_idf))
```

```
# A tibble: 119,039 x 6
```

```

  title                                bigram      n      tf    idf  tf_idf
  <chr>                                <chr> <int>  <dbl> <dbl>  <dbl>
1 The Strange Case of Dr. Jekyll and Mr. Hy~ mr ut~    69 2.92e-3  1.10 3.21e-3
2 The Strange Case of Dr. Jekyll and Mr. Hy~ the l~    61 2.58e-3  1.10 2.84e-3
3 Dracula                             van h~   282 1.88e-3  1.10 2.07e-3
4 The Strange Case of Dr. Jekyll and Mr. Hy~ mr hy~    29 1.23e-3  1.10 1.35e-3
5 The Strange Case of Dr. Jekyll and Mr. Hy~ dr je~    23 9.74e-4  1.10 1.07e-3
6 The Strange Case of Dr. Jekyll and Mr. Hy~ henry~    22 9.32e-4  1.10 1.02e-3

```



```

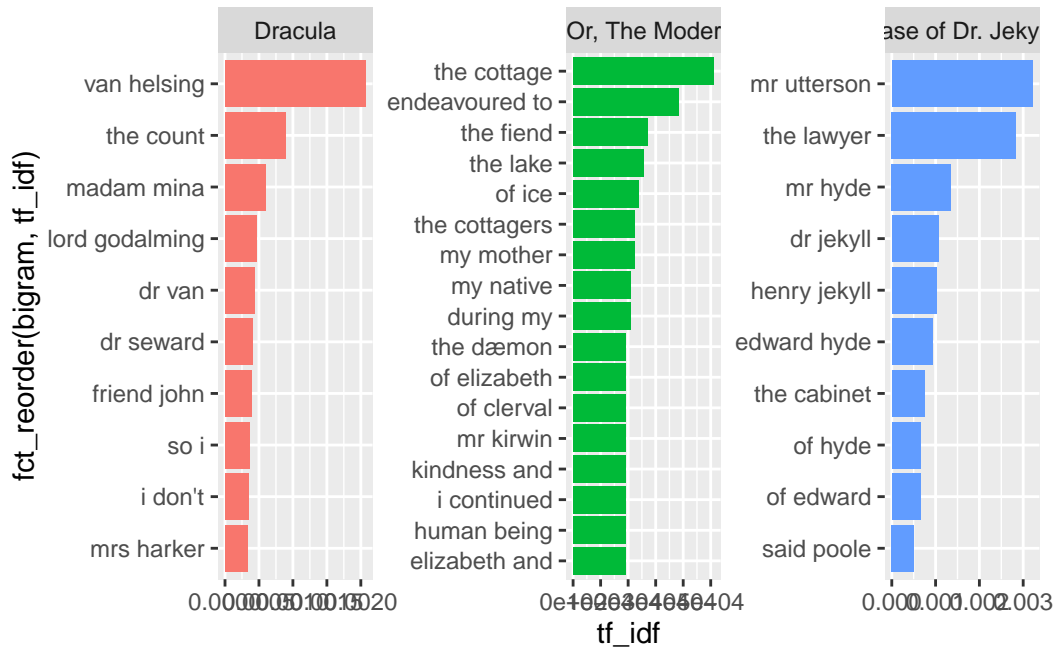
7 The Strange Case of Dr. Jekyll and Mr. Hy~ edwar~ 20 8.47e-4 1.10 9.30e-4
8 Dracula the c~ 121 8.08e-4 1.10 8.88e-4
9 The Strange Case of Dr. Jekyll and Mr. Hy~ the c~ 16 6.78e-4 1.10 7.44e-4
10 The Strange Case of Dr. Jekyll and Mr. Hy~ of ed~ 14 5.93e-4 1.10 6.51e-4
# i 119,029 more rows

```

```

bigram_tf_idf |>
  group_by(title) |>
  arrange(desc(tf_idf)) |>
  slice_max(tf_idf, n = 10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(bigram, tf_idf), y = tf_idf, fill = title)) +
    geom_col(show.legend = FALSE) +
    coord_flip() +
    facet_wrap(~title, scales = "free")

```



Sentence context using bigrams

Bigrams can also help us dive deeper into sentiment analysis. For example, even though “happy” carries positive sentiment, but when preceded by “not” as in this sentence: “I am

not happy with you!" it conveys negative sentiment. Context can matter as much as mere presence!

Let's see which words associated with an afinn sentiment are most frequently preceded by "not":

```
afinn <- get_sentiments("afinn")

bigrams_separated <- book_twowords |>
  separate(bigram, c("word1", "word2"), sep = " ") |>
  count(word1, word2, sort = TRUE) |>
  filter(!is.na(word1) & !is.na(word2))

bigrams_separated |> filter(word1 == "not")
```

```
# A tibble: 582 x 3
  word1 word2      n
  <chr> <chr> <int>
1 not   be      77
2 not   to      74
3 not   know     62
4 not   so       38
5 not   have     36
6 not   a        35
7 not   yet      34
8 not   the      31
9 not   for       29
10 not  been      26
# i 572 more rows
```

```
not_words <- bigrams_separated |>
  filter(word1 == "not") |>
  inner_join(afinn, by = c(word2 = "word")) |>
  arrange(desc(n))

not_words
```

```
# A tibble: 123 x 4
  word1 word2      n value
  <chr> <chr> <int> <dbl>
1 not   like     19      2
```

```

2 not    want      14     1
3 not    fear      13    -2
4 not    help      11     2
5 not    wish       9     1
6 not    afraid     7    -2
7 not    care       7     2
8 not    fail       7    -2
9 not    leave      7    -1
10 not   despair    6    -3
# i 113 more rows

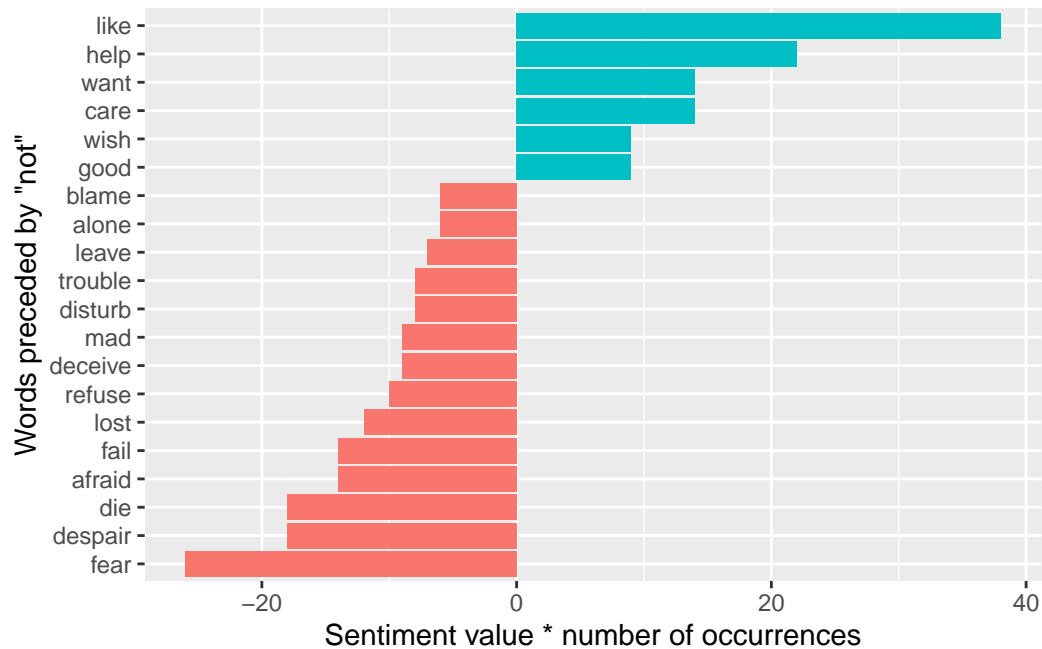
```

We could then ask which words contributed the most in the “wrong” direction. One approach is to multiply their value by the number of times they appear (so that a word with a value of +3 occurring 10 times has as much impact as a word with a sentiment value of +1 occurring 30 times).

```

not_words |>
  mutate(contribution = n * value) |>
  arrange(desc(abs(contribution))) |>
  head(20) |>
  mutate(word2 = reorder(word2, contribution)) |>
  ggplot(aes(n * value, word2, fill = n * value > 0)) +
  geom_col(show.legend = FALSE) +
  labs(x = "Sentiment value * number of occurrences",
       y = "Words preceded by \"not\"")

```



With this approach, we could expand our list of negation words, and then possibly even adjust afinn totals to reflect context!

```
# An example of expanding the list of negation words
negation_words <- c("not", "no", "never", "without")

negated_words <- bigrams_separated |>
  filter(word1 %in% negation_words) |>
  inner_join(afinn, by = c(word2 = "word")) |>
  arrange(desc(n))

negated_words
```

```
# A tibble: 232 x 4
  word1 word2      n value
  <chr> <chr>   <int> <dbl>
1 not   like     19      2
2 not   want     14      1
3 not   fear     13     -2
4 no    matter    12      1
5 no    doubt     11     -1
6 no    no        11     -1
```

```

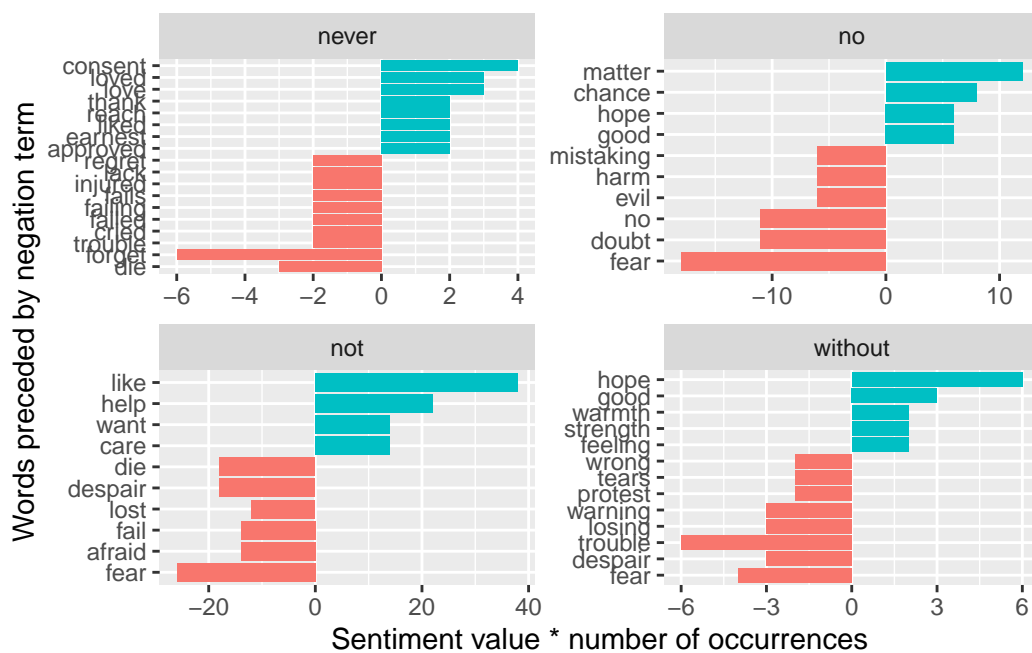
7 not    help      11     2
8 no     fear      9    -2
9 not    wish      9     1
10 not   afraid    7    -2
# i 222 more rows

```

```

negated_words |>
  mutate(contribution = n * value) |>
  arrange(desc(abs(contribution))) |>
  group_by(word1) |>
  slice_max(abs(contribution), n = 10) |>
  ungroup() |>
  mutate(word2 = reorder(word2, contribution)) |>
  ggplot(aes(n * value, word2, fill = n * value > 0)) +
    geom_col(show.legend = FALSE) +
    facet_wrap(~ word1, scales = "free") +
    labs(x = "Sentiment value * number of occurrences",
         y = "Words preceded by negation term")

```



Creating a network graph

If we are interested in visualizing all relationships among words or bigrams, we can arrange the words into a network, which is a combination of connected nodes. A network graph has three elements:

- from: the node an edge is coming from
- to: the node an edge is going towards
- weight: A numeric value associated with each edge

The `igraph` package has many powerful functions for manipulating and analyzing networks. One way to create an `igraph` object from tidy data is the `graph_from_data_frame()` function. Let's see how it works using Frankenstein:

```
library(igraph)

# filter for only relatively common combinations
bigram_graph <- bigrams_filtered |>
  filter(n > 10) |>
  graph_from_data_frame()

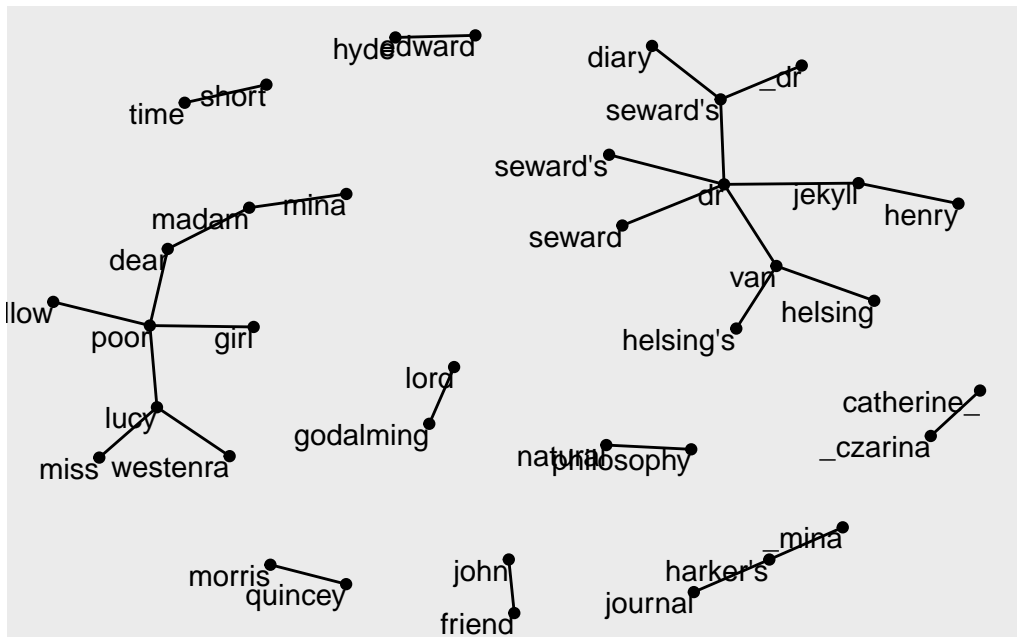
bigram_graph
```

```
IGRAPH 6177fe9 DN-- 37 27 --
+ attr: name (v/c), n (e/n)
+ edges from 6177fe9 (vertex names):
[1] van      ->helsing      madam    ->mina      lord     ->godalming
[4] dr       ->van          dr       ->seward    friend   ->john
[7] seward's->diary      poor     ->dear     harker's->journal
[10] _dr      ->seward's    dear     ->madam    miss     ->lucy
[13] dr       ->jekyll      henry    ->jekyll    poor     ->lucy
[16] quincey ->morris     edward   ->hyde     dr       ->seward's
[19] van      ->helsing's  _czarina->catherine_ poor     ->fellow
[22] _mina    ->harker's  poor     ->girl     dr       ->seward's
+ ... omitted several edges
```

```
# Use ggraph to convert into a network plot
library(ggraph)
set.seed(2017)


ggraph(bigram_graph, layout = "fr") +
  geom_edge_link() +
```

```
geom_node_point() +
geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```



```
# polish the graph
set.seed(2020)
a <- grid::arrow(type = "closed", length = unit(.15, "inches"))

ggraph(bigram_graph, layout = "fr") +
  geom_edge_link(aes(edge_alpha = n), show.legend = FALSE,
    arrow = a, end_cap = circle(.07, 'inches')) +
  geom_node_point(color = "lightblue", size = 5) +
  geom_node_text(aes(label = name), vjust = 1, hjust = 1) +
  theme_void()
```



Correlating pairs of words

Tokenizing by n-gram is a useful way to explore pairs of adjacent words. However, we may also be interested in words that tend to co-occur within particular documents or particular chapters, even if they don't occur next to each other. Following Section 4.2 of Text Mining with R, we will use the `widyr` package.

Consider the book “Frankenstein” divided into 10-line sections. We may be interested in what words tend to appear within the same section.

```
frankenstein_section_words <- frankenstein |>
  select(-gutenberg_id) |>
  mutate(section = row_number() %/% 10) |>
  filter(section > 0) |>
  unnest_tokens(word, text) |>
  filter(!word %in% stop_words$word,
         !is.na(word))

frankenstein_section_words
```

A tibble: 27,313 x 3

title

section word


```

      <chr>                                <dbl> <chr>
1 Frankenstein; Or, The Modern Prometheus    1 letter
2 Frankenstein; Or, The Modern Prometheus    1 1
3 Frankenstein; Or, The Modern Prometheus    1 letter
4 Frankenstein; Or, The Modern Prometheus    1 2
5 Frankenstein; Or, The Modern Prometheus    1 letter
6 Frankenstein; Or, The Modern Prometheus    1 3
7 Frankenstein; Or, The Modern Prometheus    1 letter
8 Frankenstein; Or, The Modern Prometheus    1 4
9 Frankenstein; Or, The Modern Prometheus    1 chapter
10 Frankenstein; Or, The Modern Prometheus    1 1
# i 27,303 more rows

```

```

# count words co-occurring within sections
library(widyr)
word_pairs <- frankenstein_section_words |>
  pairwise_count(word, section, sort = TRUE)

word_pairs

```

```

# A tibble: 856,676 x 3
   item1      item2      n
  <chr>    <chr>    <dbl>
1 elizabeth father     20
2 father   elizabeth    20
3 life     death      19
4 death    life       19
5 eyes     life       18
6 justine  poor        18
7 life     eyes        18
8 poor     justine      18
9 elizabeth dear      17
10 native  country      17
# i 856,666 more rows

```

```

# What words occur most often with "life"?
word_pairs |>
  filter(item1 == "life")

```

```

# A tibble: 2,330 x 3

```

```

  item1 item2      n
  <chr> <chr>   <dbl>
1 life  death    19
2 life  eyes     18
3 life  friend    16
4 life  father    16
5 life  mind      14
6 life  day       13
7 life  feelings  13
8 life  found     13
9 life  time      12
10 life passed    12
# i 2,320 more rows

```

We can quantify pairwise correlation using the Phi coefficient (which simplifies to the Pearson correlation coefficient with numeric data). The Phi coefficient measures how often two words appear together relative to how often they appear separately (so we don't just pick up the most common words).

```

# we need to filter for at least relatively common words first
word_cors <- frankenstein_section_words |>
  group_by(word) |>
  filter(n() >= 10) |>
  pairwise_cor(word, section, sort = TRUE)

word_cors

```

```

# A tibble: 406,406 x 3
  item1      item2 correlation
  <chr>    <chr>         <dbl>
1 philosophy natural      0.703
2 natural  philosophy    0.703
3 thou     thy          0.550
4 thy      thou          0.550
5 understood language    0.499
6 language understood    0.499
7 felix    agatha       0.470
8 agatha   felix         0.470
9 creatures fellow      0.465
10 fellow  creatures     0.465
# i 406,396 more rows

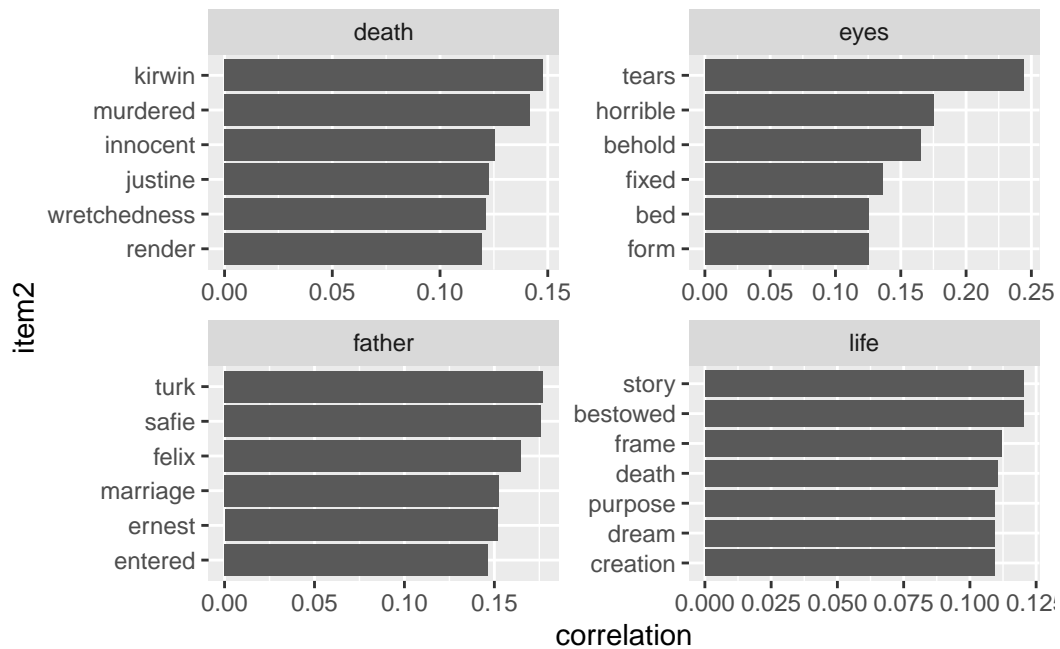
```

```
# What words are most correlated with "life"?
word_cors |>
  filter(item1 == "life")
```

```
# A tibble: 637 x 3
  item1 item2      correlation
  <chr> <chr>         <dbl>
1 life  story          0.120
2 life  bestowed        0.120
3 life  frame            0.112
4 life  death            0.111
5 life  purpose          0.109
6 life  dream            0.109
7 life  creation          0.109
8 life  deprived          0.108
9 life  hideous           0.108
10 life  money             0.105
# i 627 more rows
```

Plot words most associated with a set of interesting words:

```
word_cors |>
  filter(item1 %in% c("life", "death", "father", "eyes")) |>
  group_by(item1) |>
  slice_max(correlation, n = 6) |>
  ungroup() |>
  mutate(item2 = reorder(item2, correlation)) |>
  ggplot(aes(item2, correlation)) +
    geom_bar(stat = "identity") +
    facet_wrap(~ item1, scales = "free") +
    coord_flip()
```



Finally, create a network graph to visualize the correlations and clusters of words that were found by the widyr package

```
set.seed(2016)

word_cors |>
  filter(correlation > .25) |>
  graph_from_data_frame() |>
  ggraph(layout = "fr") +
    geom_edge_link(aes(edge_alpha = correlation), show.legend = FALSE) +
    geom_node_point(color = "lightblue", size = 5) +
    geom_node_text(aes(label = name), repel = TRUE) +
    theme_void()
```


Every document is a mixture of topics. We imagine that each document may contain words from several topics in particular proportions. For example, in a two-topic model we could say “Document 1 is 90% topic A and 10% topic B, while Document 2 is 30% topic A and 70% topic B.”

Every topic is a mixture of words. For example, we could imagine a two-topic model of American news, with one topic for “politics” and one for “entertainment.” The most common words in the politics topic might be “President”, “Congress”, and “government”, while the entertainment topic may be made up of words such as “movies”, “television”, and “actor”. Importantly, words can be shared between topics; a word like “budget” might appear in both equally.

LDA is a mathematical method for estimating both of these at the same time: finding the mixture of words that is associated with each topic, while also determining the mixture of topics that describes each document.

In order to implement LDA on our three books, we need to first “cast” our tidy data as a document-term matrix (DTM) where:

- each row represents one document (such as a book or article),
- each column represents one term, and
- each value (typically) contains the number of appearances of that term in that document.

From Section 5.2 of Text Mining with R:

Since most pairings of document and term do not occur (they have the value zero), DTMs are usually implemented as sparse matrices. These objects can be treated as though they were matrices (for example, accessing particular rows and columns), but are stored in a more efficient format.

DTM objects cannot be used directly with tidy tools, just as tidy data frames cannot be used as input for most text mining packages. Thus, the `tidytext` package provides two verbs (`tidy` and `cast`) that convert between the two formats.

A DTM is typically comparable to a tidy data frame after a `count` or a `group_by/summarize` that contains counts or another statistic for each combination of a term and document.

```
# cast the collection of 3 works as a document-term matrix
library(tm)
```

Loading required package: NLP

Attaching package: 'NLP'

The following object is masked from 'package:ggplot2':

annotate

```
three_books_dtm <- book_word_count |>
  filter(!word %in% stop_words$word,
         !is.na(word)) |>
  cast_dtm(title, word, n)

# set a seed so that the output of the model is predictable
library(topicmodels)
three_books_lda <- LDA(three_books_dtm, k = 2, control = list(seed = 1234))
three_books_lda
```

A LDA_VEM topic model with 2 topics.

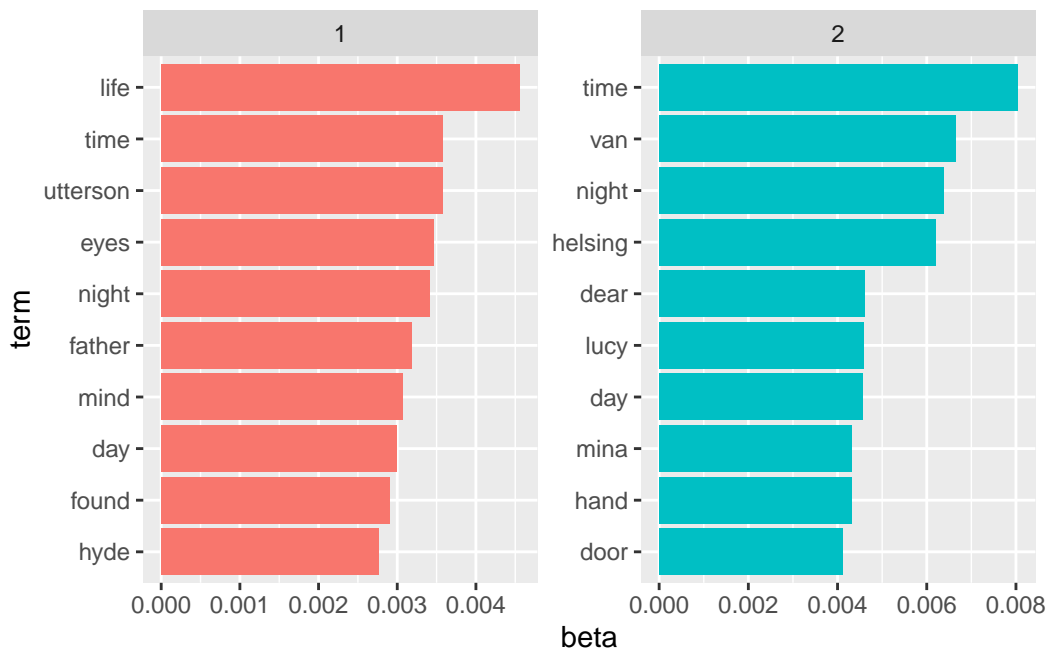
After fitting our LDA model, we will first focus on the beta variable, which is the probability of a word being generated by a specific topic. Then we'll turn to the gamma variable, which are the per-document per-topic probabilities, or the proportion of words from a document generated by a specific topic.

```
three_books_topics <- tidy(three_books_lda, matrix = "beta")
three_books_topics
```

```
# A tibble: 25,968 x 3
  topic term      beta
  <int> <chr>    <dbl>
1     1 time  3.58e- 3
2     2 time  8.03e- 3
3     1 van   1.26e-14
4     2 van   6.65e- 3
5     1 night 3.41e- 3
6     2 night 6.39e- 3
7     1 helsing 1.91e-14
8     2 helsing 6.20e- 3
9     1 dear  2.18e- 3
10    2 dear  4.61e- 3
# i 25,958 more rows
```

```
# Find the most common words within each topic
three_books_top_terms <- three_books_topics |>
  group_by(topic) |>
  slice_max(beta, n = 10) |>
  ungroup() |>
  arrange(topic, -beta)

three_books_top_terms |>
  mutate(term = reorder_within(term, beta, topic)) |>
  ggplot(aes(beta, term, fill = factor(topic))) +
    geom_col(show.legend = FALSE) +
    facet_wrap(~ topic, scales = "free") +
    scale_y_reordered()
```



```
# This would be much cooler with more documents and if we were able
# to anti_join to remove proper nouns

# Find words with greatest difference between two topics, using log ratio
beta_wide <- three_books_topics |>
  mutate(topic = paste0("topic", topic)) |>
  pivot_wider(names_from = topic, values_from = beta) |>
```

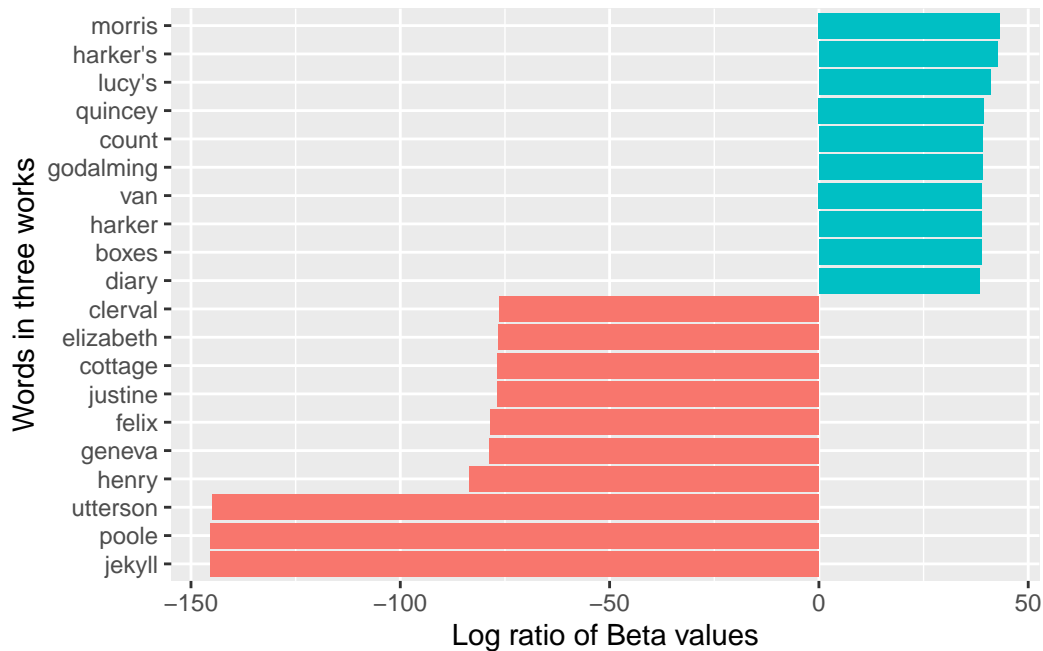


```
filter(topic1 > .001 | topic2 > .001) |>
mutate(log_ratio = log2(topic2 / topic1))
```

```
beta_wide
```

```
# A tibble: 196 x 4
  term      topic1 topic2 log_ratio
  <chr>    <dbl>   <dbl>   <dbl>
1 time    3.58e- 3 0.00803    1.17
2 van     1.26e-14 0.00665   38.9
3 night   3.41e- 3 0.00639    0.906
4 helsing 1.91e-14 0.00620   38.2
5 dear    2.18e- 3 0.00461    1.08
6 lucy     5.62e-14 0.00459   36.3
7 day     2.99e- 3 0.00455    0.607
8 hand    1.98e- 3 0.00433    1.12
9 mina    2.40e-14 0.00433   37.4
10 door    2.12e- 3 0.00412    0.956
# i 186 more rows
```

```
beta_wide |>
  arrange(desc(abs(log_ratio))) |>
  slice_max(abs(log_ratio), n = 20) |>
  mutate(term = reorder(term, log_ratio)) |>
  ggplot(aes(log_ratio, term, fill = log_ratio > 0)) +
    geom_col(show.legend = FALSE) +
    labs(x = "Log ratio of Beta values",
         y = "Words in three works")
```



```
# find the gamma variable for each document and topic
three_books_documents <- tidy(three_books_lda, matrix = "gamma")
three_books_documents
```

```
# A tibble: 6 x 3
  document                                topic    gamma
  <chr>                                <int>    <dbl>
1 Dracula                               1 0.000158
2 The Strange Case of Dr. Jekyll and Mr. Hyde 1 1.00
3 Frankenstein; Or, The Modern Prometheus 1 1.00
4 Dracula                               2 1.00
5 The Strange Case of Dr. Jekyll and Mr. Hyde 2 0.00000613
6 Frankenstein; Or, The Modern Prometheus 2 0.00000190
```

```
# Dracula = Topic 2; other two books = Topic 1!
```

On Your Own: Harry Potter

The `potter_untidy` dataset includes the text of 7 books of the Harry Potter series by J.K. Rowling. For a brief overview of the books (or movies), see this quote from Wikipedia:

Harry Potter is a series of seven fantasy novels written by British author J. K. Rowling. The novels chronicle the lives of a young wizard, Harry Potter, and his friends Hermione Granger and Ron Weasley, all of whom are students at Hogwarts School of Witchcraft and Wizardry. The main story arc concerns Harry's conflict with Lord Voldemort, a dark wizard who intends to become immortal, overthrow the wizard governing body known as the Ministry of Magic, and subjugate all wizards and Muggles (non-magical people).

A few analyses from SDS 164:

```
# 10 most common words in each book, excluding stop words
potter_tidy |>
  count(title, word) |>
  anti_join(stop_words) |>
  group_by(title) |>
  slice_max(n, n = 10) |>
  mutate(rank = 1:10) |>
  select(-n) |>
  pivot_wider (names_from = title, values_from = word) |>
  print(width = Inf)
```

Joining with `by = join_by(word)`

```
# A tibble: 10 x 8
  rank `Sorcerer's Stone` `Chamber of Secrets` `Prisoner of Azkaban`
  <int> <chr>              <chr>              <chr>
1     1 harry             harry             harry
2     2 ron              ron              ron
3     3 hagrid           hermione          hermione
4     4 hermione         malfoy           professor
5     5 professor        lockhart         lupin
6     6 looked           professor        black
7     7 snape            weasley          looked
8     8 dumbledore       looked          hagrid
9     9 uncle            time             snape
10    10 time            eyes             harry's
  `Goblet of Fire` `Order of the Phoenix` `Half-Blood Prince` `Deathly Hallows`
  <chr>            <chr>            <chr>            <chr>
1 harry           harry           harry           harry
2 ron            hermione        dumbledore      hermione
```

3	hermione	ron	ron	ron
4	dumbledore	sirius	hermione	wand
5	looked	professor	looked	dumbledore
6	weasley	dumbledore	slughorn	looked
7	hagrid	looked	snape	voldemort
8	eyes	umbridge	malfoy	eyes
9	moody	weasley	time	death
10	professor	voice	professor	time

```
# Repeat above after removing character first and last names
potter_tidy |>
  count(title, word) |>
  anti_join(stop_words) |>
  anti_join(potter_names, join_by(word == firstname)) |>
  anti_join(potter_names, join_by(word == lastname)) |>
  group_by(title) |>
  slice_max(n, n = 10, with_ties = FALSE) |>
  mutate(rank = 1:10) |>
  select(-n) |>
  pivot_wider (names_from = title, values_from = word) |>
  print(width = Inf)
```

Joining with `by = join_by(word)`

```
# A tibble: 10 x 8
  rank `Sorcerer's Stone` `Chamber of Secrets` `Prisoner of Azkaban`
  <int> <chr>              <chr>              <chr>
1     1 professor         professor         professor
2     2 looked           looked           looked
3     3 uncle            time             harry's
4     4 time             eyes            eyes
5     5 harry's          harry's          time
6     6 door             doobby          door
7     7 eyes            door            head
8     8 yeh              head            voice
9     9 head             voice          heard
10    10 told            school          hand
  `Goblet of Fire` `Order of the Phoenix` `Half-Blood Prince` `Deathly Hallows`
  <chr>            <chr>            <chr>            <chr>
1 looked          professor        looked          wand
2 eyes            looked         time            looked
```

3	professor	voice	professor	eyes
4	crouch	time	hand	death
5	time	door	eyes	time
6	wand	head	voice	voice
7	voice	harry's	dark	harry's
8	head	eyes	wand	door
9	told	wand	door	hand
10	harry's	hand	head	head

```
# still get "harry's" and "professor" but otherwise looks good
```

```
# top 10 names in each book (after excluding "the")
```

```
potter_tidy |>
  count(title, word) |>
  semi_join(potter_names, join_by(word == firstname)) |>
  filter(word != "the") |> # ADD for #6
  group_by(title) |>
  slice_max(n, n = 10, with_ties = FALSE) |>
  mutate(rank = 1:10) |>
  select(-n) |>
  pivot_wider(names_from = title, values_from = word) |>
  print(width = Inf)
```

```
# A tibble: 10 x 8
```

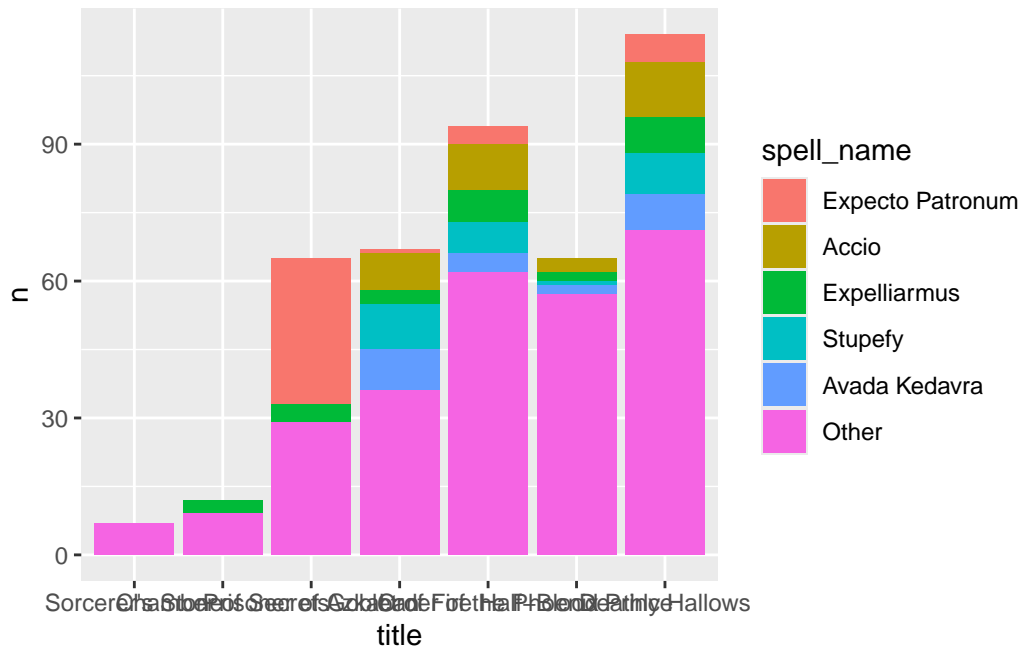
	rank	`Sorcerer's Stone`	`Chamber of Secrets`	`Prisoner of Azkaban`	
	<int>	<chr>	<chr>	<chr>	
1	1	harry	harry	harry	
2	2	ron	ron	ron	
3	3	hermione	hermione	hermione	
4	4	dudley	fred	sirius	
5	5	vernon	ginny	neville	
6	6	neville	sir	madam	
7	7	great	george	great	
8	8	petunia	great	fred	
9	9	nearly	percy	vernon	
10	10	madam	nearly	percy	
		`Goblet of Fire`	`Order of the Phoenix`	`Half-Blood Prince`	`Deathly Hallows`
		<chr>	<chr>	<chr>	<chr>
1		harry	harry	harry	harry
2		ron	hermione	ron	hermione
3		hermione	ron	hermione	ron

4	cedric	sirius	ginny	great
5	sirius	fred	great	lord
6	fred	george	sir	luna
7	great	neville	lord	bill
8	george	ginny	fred	ginny
9	percy	great	tom	albus
10	rita	luna	draco	fred

```
# spell statistics by book
potter_tidy |>
  left_join(potter_spells, join_by(word == first_word)) |>
  group_by(title) |>
  summarize(num_spells_cast = sum(!is.na(spell_name)),
            spells_per_10kwords = mean(!is.na(spell_name)) * 10000,
            num_unique_spells = n_distinct(spell_name) - 1) # Why -1??
```

```
# A tibble: 7 x 4
  title                num_spells_cast spells_per_10kwords num_unique_spells
  <fct>                <int>          <dbl>          <dbl>
1 Sorcerer's Stone      7            0.899            4
2 Chamber of Secrets   12            1.41            9
3 Prisoner of Azkaban   65            6.17           14
4 Goblet of Fire        67            3.49           27
5 Order of the Phoenix  94            3.63           28
6 Half-Blood Prince     65            3.79           24
7 Deathly Hallows     114            5.77           34
```

```
# plot of top spells by book
potter_tidy |>
  left_join(potter_spells, join_by(word == first_word)) |>
  drop_na(spell_name) |>
  mutate(spell_name = fct_infreq(spell_name),
         spell_name = fct_lump_n(spell_name, n = 5)) |>
  count(title, spell_name) |>
  ggplot() +
  geom_col(aes(x = title, y = n, fill = spell_name), position = "stack")
```



New stuff!


1. What words contribute the most to negative and positive sentiment scores? Show a faceted bar plot of the top 10 negative and the top 10 positive words (according to the “bing” lexicon) across the entire series.

```
bing_sentiments <- get_sentiments(lexicon = "bing")

potter_tidy |>
  inner_join(bing_sentiments) |>
  count(sentiment, word, sort = TRUE) |>
  group_by(sentiment) |>
  slice_max(n, n = 10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(word, n), y = n, fill = sentiment)) +
    geom_col() +
    coord_flip() +
    facet_wrap(~ sentiment, scales = "free")
```

Joining with `by = join_by(word)`

```
Warning in inner_join(potter_tidy, bing_sentiments): Detected an unexpected many-to-many relationship.
i Row 41432 of `x` matches multiple rows in `y`.
i Row 2698 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to silence this warning.
```



- Find a list of the top 10 words associated with “fear” and with “trust” (according to the “nrc” lexicon) across the entire series.

```
potter_tidy |>
  inner_join(nrc_sentiments) |>
  filter(sentiment == c("fear", "trust")) |>
  count(sentiment, word, sort = TRUE) |>
  group_by(sentiment) |> #I'm not sure if you mean the 10 ten words associated with fear a
  slice_max(n, n = 10, with_ties = FALSE) |>
  print(n = Inf)
```

Joining with `by = join_by(word)`

```
Warning in inner_join(potter_tidy, nrc_sentiments): Detected an unexpected many-to-many relationship.
i Row 15 of `x` matches multiple rows in `y`.
```



```
i Row 11469 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
  "many-to-many"` to silence this warning.
```

```
Warning: There was 1 warning in `filter()`.
i In argument: `sentiment == c("fear", "trust")`.
Caused by warning in `sentiment == c("fear", "trust")`:
! longer object length is not a multiple of shorter object length
```

```
# A tibble: 20 x 3
# Groups:   sentiment [2]
  sentiment word      n
  <chr>      <chr> <int>
1 fear      death    377
2 fear      feeling   198
3 fear      fire     185
4 fear      crouch    152
5 fear      mad       152
6 fear      kill      143
7 fear      scar      142
8 fear      shaking   135
9 fear      darkness  131
10 fear     bad      126
11 trust     professor  968
12 trust     good      544
13 trust     found     304
14 trust     ministry  299
15 trust     school    290
16 trust     sir       208
17 trust     top       204
18 trust     lord      196
19 trust     ground    194
20 trust     feeling   193
```

3. Make a wordcloud for the entire series after removing stop words using the “smart” source.


```
smart_stopwords <- get_stopwords(source = "smart")

hp_words <- potter_tidy |>
  anti_join(smart_stopwords) |>
  count(word) |>
```

```
filter(word != "NA") |>
  arrange(desc(n))
```

Joining with ``by = join_by(word)``

```
wordcloud(  
  words = hp_words$word,  
  freq = hp_words$n,  
  max.words = 100,  
  random.order = FALSE  
)
```



4. Create a wordcloud with the top 20 negative words and the top 20 positive words in the Harry Potter series according to the bing lexicon. The words should be sized by their respective counts and colored based on whether their sentiment is positive or negative. (Feel free to be resourceful and creative to color words by a third variable!)

```
pos_neg <- potter_tidy |>
  inner_join(get_sentiments("bing")) |>
  count(sentiment, word, sort = TRUE) |>
```

```
group_by(sentiment) |>
top_n(20) |>
ungroup()
```

Joining with `by = join_by(word)`

```
Warning in inner_join(potter_tidy, get_sentiments("bing")): Detected an unexpected many-to-many
i Row 41432 of `x` matches multiple rows in `y`.
i Row 2698 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
  "many-to-many"` to silence this warning.
```

Selecting by n

```
wordcloud(
  word = pos_neg$word,
  freq = pos_neg$n,
  random.order = FALSE,
  rot.per = 0,
  ordered.colors = TRUE,
  colors = brewer.pal(6, "Dark2")[factor(pos_neg$sentiment)]
)
```




6. Repeat (5) using a faceted scatterplot to show the average sentiment score according to the “afinn” lexicon for each chapter. (Hint: use `mutate(chapter_factor = factor(chapter))` to treat chapter as a factor variable.)


```
potter_tidy |>
  inner_join(get_sentiments("afinn")) |>
  mutate(chapter_factor = factor(chapter)) |>
  group_by(title, chapter) |>
  mutate(sentiment = mean(value)) |>
  ggplot(aes(x = chapter, y = sentiment, fill = title, color = title)) +
    geom_point(show.legend = FALSE) +
    facet_wrap(~title, ncol = 2, scales = "free_x")
```

Joining with ``by = join_by(word)``



7. Make a faceted bar plot showing the top 10 words that distinguish each book according to the tf-idf statistic.

```
potter_tidy |>
  count(word, title, sort = TRUE) |>
  bind_tf_idf(word, title, n) |>
  group_by(title) |>
  arrange(desc(tf_idf)) |>
  slice_max(tf_idf, n = 10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(word, tf_idf), y = tf_idf, fill = title)) +
    geom_col(show.legend = FALSE) +
    coord_flip() +
    facet_wrap(~title, scales = "free")
```



8. Repeat (7) to show the top 10 2-word combinations that distinguish each book.

```
hp_twowords <- potter_untidy |>
  unnest_tokens(bigram, text, token = "ngrams", n = 2) |>
  filter(bigram != "NA") |>
  separate(bigram, c("word1", "word2"), sep = " ") |>
  filter(!word1 %in% stop_words$word,
         !word2 %in% stop_words$word) |>
  filter(!is.na(word1) & !is.na(word2))
```

hp_twowords

A tibble: 137,802 x 5

	title	chapter	book_num	word1	word2
	<fct>	<dbl>	<dbl>	<chr>	<chr>
1	Sorcerer's Stone	1	1	privet	drive
2	Sorcerer's Stone	1	1	perfectly	normal
3	Sorcerer's Stone	1	1	firm	called
4	Sorcerer's Stone	1	1	called	grunnings
5	Sorcerer's Stone	1	1	usual	amount
6	Sorcerer's Stone	1	1	time	craning
7	Sorcerer's Stone	1	1	garden	fences

```

8 Sorcerer's Stone      1      1 fences    spying
9 Sorcerer's Stone      1      1 son      called
10 Sorcerer's Stone     1      1 called   dudley
# i 137,792 more rows

```

```

hp_bigrams_united <- hp_twowords |>
  unite(bigram, word1, word2, sep = " ")

```

```
hp_bigrams_united
```

```

# A tibble: 137,802 x 4
  title          chapter book_num bigram
<fct>          <dbl>    <dbl> <chr>
1 Sorcerer's Stone      1      1 privet drive
2 Sorcerer's Stone      1      1 perfectly normal
3 Sorcerer's Stone      1      1 firm called
4 Sorcerer's Stone      1      1 called grunnings
5 Sorcerer's Stone      1      1 usual amount
6 Sorcerer's Stone      1      1 time craning
7 Sorcerer's Stone      1      1 garden fences
8 Sorcerer's Stone      1      1 fences spying
9 Sorcerer's Stone      1      1 son called
10 Sorcerer's Stone     1      1 called dudley
# i 137,792 more rows

```

```

hp_bigram_tf_idf <- hp_bigrams_united |>
  count(title, bigram) |>
  bind_tf_idf(bigram, title, n) |>
  arrange(desc(tf_idf))

```

```
hp_bigram_tf_idf |> arrange(desc(tf_idf))
```

```

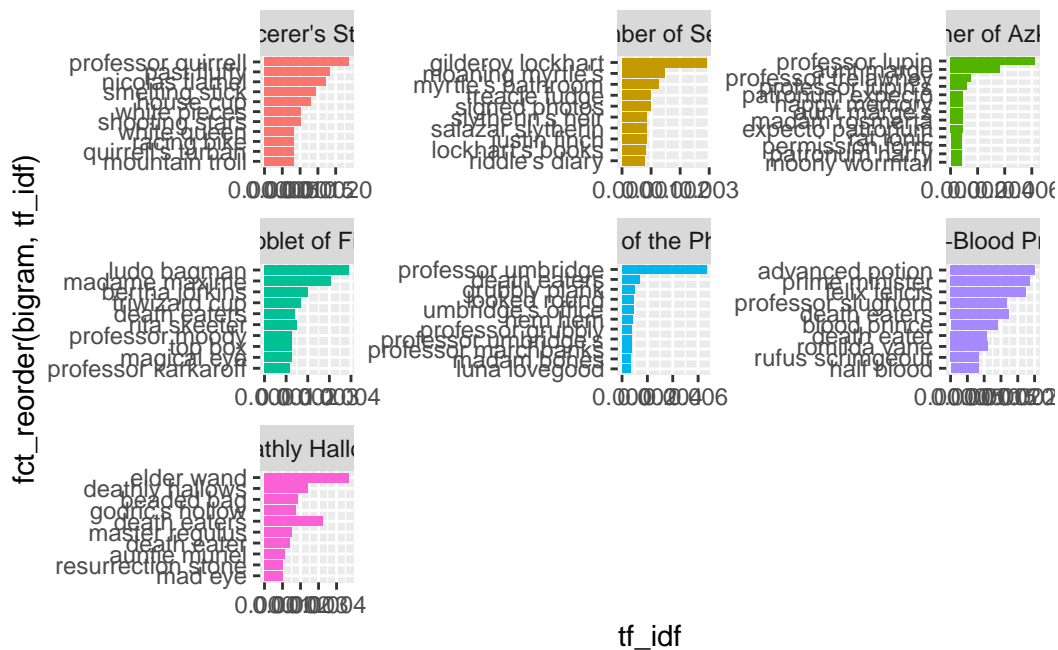
# A tibble: 107,166 x 6
  title          bigram      n    tf  idf  tf_idf
<fct>          <chr>    <int> <dbl> <dbl> <dbl>
1 Order of the Phoenix professor umbridge 173 0.00533 1.25 0.00667
2 Prisoner of Azkaban professor lupin    107 0.00738 0.847 0.00625
3 Deathly Hallows    elder wand       58 0.00241 1.95 0.00470
4 Goblet of Fire     ludo bagman      49 0.00201 1.95 0.00391
5 Prisoner of Azkaban aunt marge       42 0.00290 1.25 0.00363

```


6	Deathly Hallows	death eaters	139	0.00578	0.560	0.00324
7	Goblet of Fire	madame maxime	89	0.00365	0.847	0.00309
8	Chamber of Secrets	gilderoy lockhart	28	0.00232	1.25	0.00291
9	Half-Blood Prince	advanced potion	27	0.00129	1.95	0.00252
10	Deathly Hallows	deathly hallows	30	0.00125	1.95	0.00243

i 107,156 more rows

```
hp_bigram_tf_idf |>
  group_by(title) |>
  arrange(desc(tf_idf)) |>
  slice_max(tf_idf, n = 10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(bigram, tf_idf), y = tf_idf, fill = title)) +
    geom_col(show.legend = FALSE) +
    coord_flip() +
    facet_wrap(~title, scales = "free")
```



- Find which words contributed most in the “wrong” direction using the afinn sentiment combined with how often a word appears among all 7 books. Come up with a list of 4 negation words, and for each negation word, illustrate the words associated with the largest “wrong” contributions in a faceted bar plot.

```

afinn <- get_sentiments("afinn")

hp_twowords2 <- potter_untidy |>
  unnest_tokens(bigram, text, token = "ngrams", n = 2) |>
  filter(bigram != "NA")

hp_bigrams_separated <- hp_twowords2 |>
  separate(bigram, c("word1", "word2"), sep = " ") |>
  count(word1, word2, sort = TRUE) |>
  filter(!is.na(word1) & !is.na(word2))

hp_negation_words <- c("neither", "nor", "not", "never")

hp_negated_words <- hp_bigrams_separated |>
  filter(word1 %in% hp_negation_words) |>
  inner_join(afinn, by = c(word2 = "word")) |>
  arrange(desc(n))

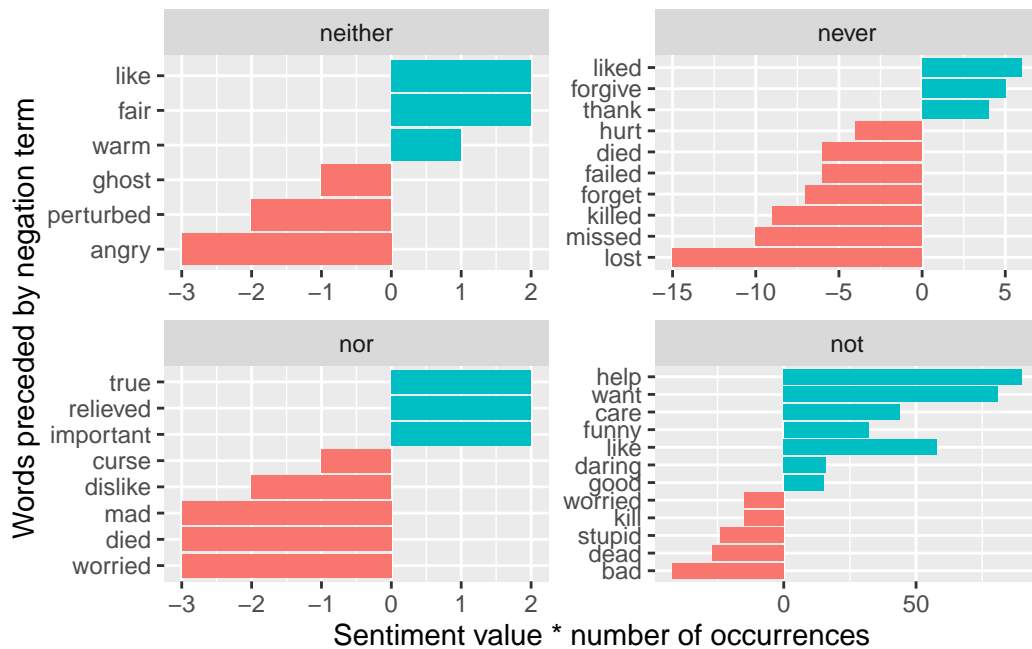
hp_negated_words

# A tibble: 273 x 4
  word1 word2      n value
  <chr> <chr>   <int> <dbl>
1 not    want      81      1
2 not    help      45      2
3 not    like      29      2
4 not    care      22      2
5 not    bad       14     -3
6 not    wish      14      1
7 not    stupid    12     -2
8 not    stop      11     -1
9 not    matter    10      1
10 not    dead       9     -3
# i 263 more rows

hp_negated_words |>
  mutate(contribution = n * value) |>
  arrange(desc(abs(contribution))) |>
  group_by(word1) |>
  slice_max(abs(contribution), n = 10) |>
  ungroup() |>

```

```
mutate(word2 = reorder(word2, contribution)) |>
ggplot(aes(n * value, word2, fill = n * value > 0)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ word1, scales = "free") +
  labs(x = "Sentiment value * number of occurrences",
       y = "Words preceded by negation term")
```



10. Select a set of 4 “interesting” terms and then use the Phi coefficient to find and plot the 6 words most correlated with each of your “interesting” words. Start by dividing `potter_tidy` into 80-word sections and then remove names and spells and stop words.

```
library(widyr)

potter_section_words <- potter_tidy |>
mutate(section = 1 + row_number() %/% 80) |>
filter(!word %in% stop_words$word,
       !is.na(word),
       !word %in% potter_names$firstname,
       !word %in% potter_names$lastname,
       !word %in% potter_spells$spell_name,
       !word %in% potter_spells$spell,
```

```
!word %in% potter_spells$first_word,
!word %in% potter_spells$second_word)
```

```
potter_section_words
```

```
# A tibble: 346,372 x 5
```

	title	chapter	book_num	word	section
	<fct>	<dbl>	<dbl>	<chr>	<dbl>
1	Sorcerer's Stone	1	1	boy	1
2	Sorcerer's Stone	1	1	lived	1
3	Sorcerer's Stone	1	1	privet	1
4	Sorcerer's Stone	1	1	drive	1
5	Sorcerer's Stone	1	1	proud	1
6	Sorcerer's Stone	1	1	perfectly	1
7	Sorcerer's Stone	1	1	normal	1
8	Sorcerer's Stone	1	1	people	1
9	Sorcerer's Stone	1	1	expect	1
10	Sorcerer's Stone	1	1	involved	1

```
# i 346,362 more rows
```

```
potter_word_cors <- potter_section_words |>
  group_by(word) |>
  filter(n() >= 10) |>
  pairwise_cor(word, section, sort = TRUE)
```

```
potter_word_cors
```

```
# A tibble: 28,222,656 x 3
```

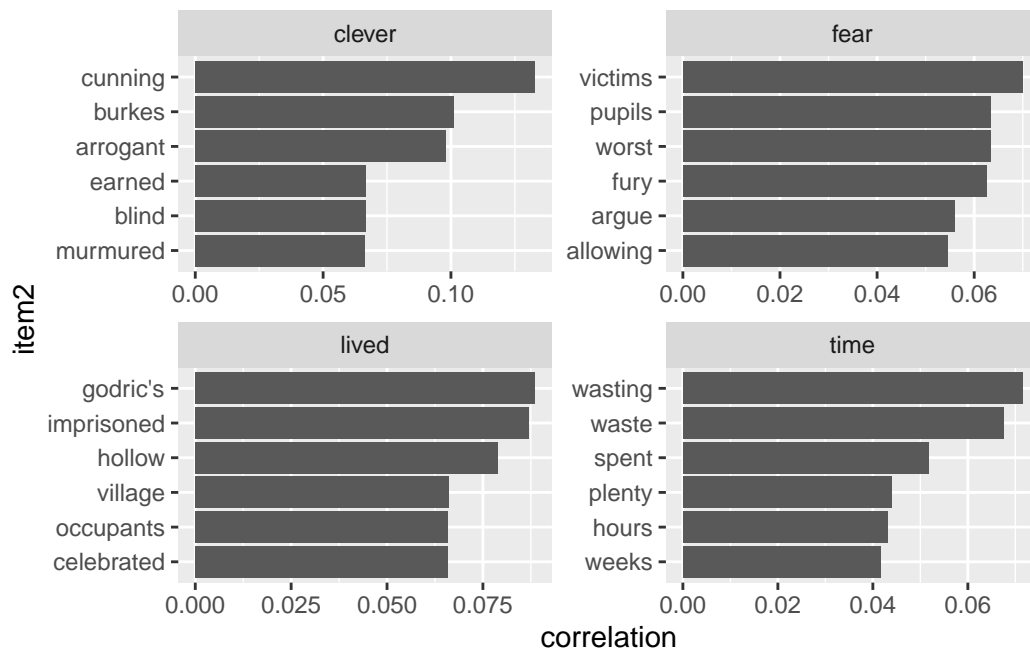
	item1	item2	correlation
	<chr>	<chr>	<dbl>
1	snare	devil's	1
2	devil's	snare	1
3	grubbly	plank	0.925
4	plank	grubbly	0.925
5	crescent	magnolia	0.904
6	magnolia	crescent	0.904
7	fletchley	finch	0.897
8	finch	fletchley	0.897
9	prophet	daily	0.852
10	daily	prophet	0.852

```
# i 28,222,646 more rows
```

```

potter_word_cors |>
  filter(item1 %in% c("clever", "lived", "time", "fear")) |>
  group_by(item1) |>
  slice_max(correlation, n = 6) |>
  ungroup() |>
  mutate(item2 = reorder(item2, correlation)) |>
  ggplot(aes(item2, correlation)) +
    geom_bar(stat = "identity") +
    facet_wrap(~ item1, scales = "free") +
    coord_flip()

```



11. Create a network graph to visualize the correlations and clusters of words that were found by the `widyr` package in (10).

```

library(igraph)
set.seed(7777)

potter_word_cors |>
  filter(correlation > .5) |>
  graph_from_data_frame() |>
  ggraph(layout = "fr") +

```



```

!word %in% potter_spells$spell_name,
!word %in% potter_spells$spell,
!word %in% potter_spells$first_word,
!word %in% potter_spells$second_word) |>
cast_dtm(title, word, n)

library(topicmodels)
potter_tidy_lda <- LDA(potter_tidy_dtm, k = 2, control = list(seed = 2003))
potter_tidy_lda

```

A LDA_VEM topic model with 2 topics.

```

potter_books_topics <- tidy(potter_tidy_lda, matrix = "beta")
three_books_topics

```

```

# A tibble: 25,968 x 3
  topic term      beta
  <int> <chr>    <dbl>
1     1 time  3.58e- 3
2     2 time  8.03e- 3
3     1 van   1.26e-14
4     2 van   6.65e- 3
5     1 night 3.41e- 3
6     2 night 6.39e- 3
7     1 helsing 1.91e-14
8     2 helsing 6.20e- 3
9     1 dear  2.18e- 3
10    2 dear  4.61e- 3
# i 25,958 more rows

```

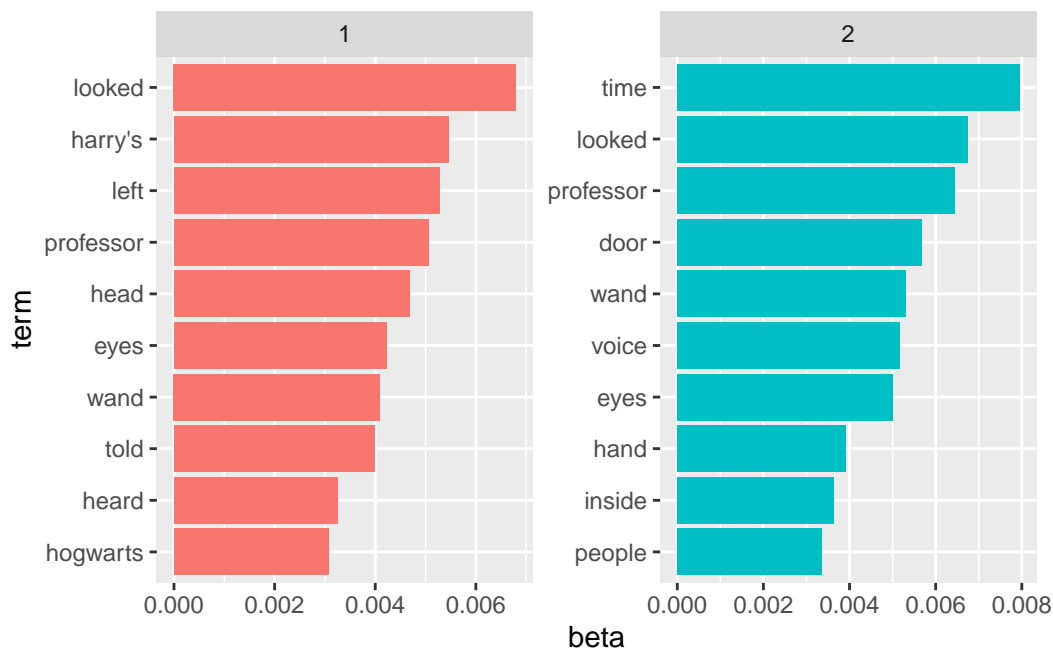
```

potter_books_top_terms <- potter_books_topics |>
  group_by(topic) |>
  slice_max(beta, n = 10) |>
  ungroup() |>
  arrange(topic, -beta)

potter_books_top_terms |>
  mutate(term = reorder_within(term, beta, topic)) |>
  ggplot(aes(beta, term, fill = factor(topic))) +

```

```
geom_col(show.legend = FALSE) +
facet_wrap(~ topic, scales = "free") +
scale_y_reordered()
```



A.) A plot to illustrate words with greatest difference between two topics, using log ratio.

```
potter_beta <- potter_books_topics |>
  mutate(topic = paste0("topic", topic)) |>
  pivot_wider(names_from = topic, values_from = beta) |>
  filter(topic1 > .001 | topic2 > .001) |>
  mutate(log_ratio = log2(topic2 / topic1))
```

potter_beta

A tibble: 191 x 4

	term	topic1	topic2	log_ratio
	<chr>	<dbl>	<dbl>	<dbl>
1	professor	0.00506	0.00645	0.351
2	wand	0.00410	0.00530	0.373
3	looked	0.00680	0.00674	-0.0124
4	voice	0.00305	0.00517	0.763


```


5 time      0.00160 0.00794    2.31
6 door      0.00163 0.00567    1.80
7 head      0.00468 0.00307   -0.609
8 harry's   0.00545 0.00266   -1.04
9 eyes      0.00422 0.00500    0.242
10 death    0.00249 0.00191   -0.385
# i 181 more rows

```

```

potter_beta |>
  arrange(desc(abs(log_ratio))) |>
  slice_max(abs(log_ratio), n = 20) |>
  mutate(term = reorder(term, log_ratio)) |>
  ggplot(aes(log_ratio, term, fill = log_ratio > 0)) +
    geom_col(show.legend = FALSE) +
    labs(x = "Log ratio of Beta values",
         y = "Words in the Harry Potter Series")

```



B.) A table with the gamma variable for each document and topic.

```

potter_documents <- tidy(potter_tidy_lda, matrix = "gamma")
potter_documents

```

```
# A tibble: 14 x 3
  document      topic gamma
  <chr>         <int> <dbl>
1 Sorcerer's Stone      1 0.506
2 Chamber of Secrets    1 0.477
3 Prisoner of Azkaban    1 0.524
4 Goblet of Fire         1 0.464
5 Order of the Phoenix   1 0.483
6 Half-Blood Prince      1 0.484
7 Deathly Hallows       1 0.427
8 Sorcerer's Stone      2 0.494
9 Chamber of Secrets    2 0.523
10 Prisoner of Azkaban   2 0.476
11 Goblet of Fire        2 0.536
12 Order of the Phoenix  2 0.517
13 Half-Blood Prince     2 0.516
14 Deathly Hallows      2 0.573
```

- Looking at the gamma variables, all of the books are pretty evenly split between topic 1 and topic 2, which isn't too surprising given that the books seem pretty similar content-wise. Looking at the words with the greatest different between the two topics a possibility for what the two topics represent is that one topic represents anticipation and success and the other represents challenges.